MemoryOS — The Persistent Memory Layer for AI Agents

The "Redis + Postgres" for AI Agent Memory — $30B+ Market Opportunity

---

Executive Summary

MemoryOS is the foundational infrastructure for persistent, intelligent memory in AI agents. Every AI agent today wakes up with amnesia — no recall of past interactions, learned preferences, or accumulated context. This is the single biggest limitation holding back the agentic AI revolution.

MemoryOS solves this with a unified memory layer that provides: **persistent storage, semantic retrieval, privacy controls, memory sharing, and cognitive optimization** — all through a simple API that any AI agent can integrate in minutes.

The thesis: Every AI agent needs memory. We become the universal memory infrastructure for the 100M+ AI agents deployed by 2028. This is the "database layer" for the agentic era.

The Problem

The AI Amnesia Crisis

The 2025-2026 AI agent explosion has exposed a critical infrastructure gap:

1. **Agents forget everything** — Every session starts fresh. Your AI assistant doesn't remember your preferences, your history, or what you discussed yesterday.
2. **Context windows are finite** — Even with 200K+ token windows, agents can't retain months or years of interaction history.
3. **Memory systems are fragmented** — Every company builds their own vector DB + retrieval stack. Massive duplication of effort.
4. **No standard memory primitives** — There's no agreed-upon format for how agents should store, retrieve, and reason about memories.
5. **Privacy is an afterthought** — User data ends up scattered across vector databases with no consent management, retention policies, or right-to-delete.
6. **No memory portability** — Switch from ChatGPT to Claude? Your "memory" doesn't transfer. Vendor lock-in through data captivity.

The Market Reality

- **OpenAI** launched "Memory" in ChatGPT — but it's basic, siloed, and platform-locked
- **Anthropic, Google, Microsoft** all building proprietary memory systems
- **Enterprise agents** (customer support, sales, operations) desperately need persistent context
- **Personal AI assistants** are useless without remembering who you are
- **Multi-agent systems** need shared memory to coordinate effectively

The market is screaming for a standard. No one has built it yet.

The Solution

MemoryOS Platform

Universal memory infrastructure for any AI agent, any platform, any use case.

Core Architecture

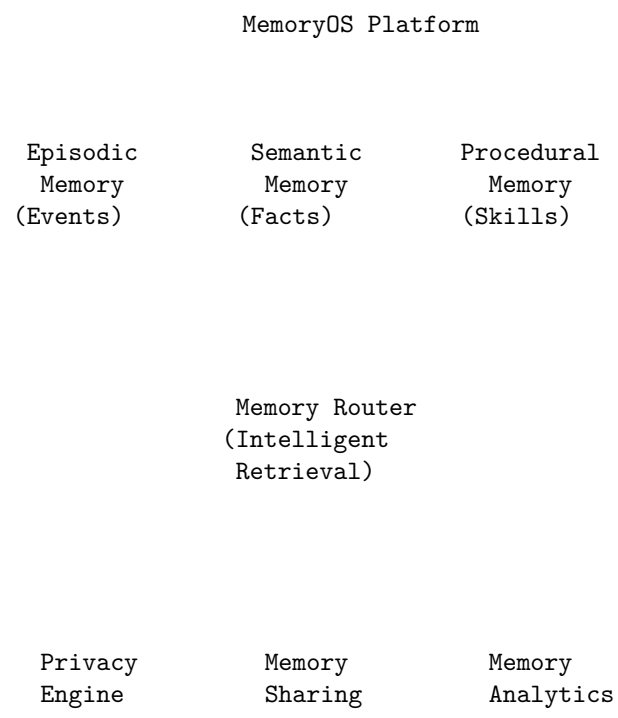

Memory Types (Inspired by Cognitive Science) **1. Episodic Memory — "What Happened"** - Stores specific events, conversations, interactions - Time-stamped, contextual, retrievable by similarity or time - "Remember when I mentioned I'm allergic to shellfish?"

2. Semantic Memory — "What I Know" - Stores facts, preferences, learned information - Structured knowledge graphs + vector embeddings - "User prefers dark mode, lives in NYC, works at Acme Corp"

3. Procedural Memory — "How To Do Things" - Stores learned workflows, patterns, skills - "When user says 'deploy', run these 5 steps in this order"

4. Working Memory — "Active Context" - Optimized short-term store for current session - Intelligent compression and summarization - Seamless overflow to long-term storage

---

Product Suite

1. MemoryOS Core (Infrastructure)

The foundational memory layer:

```
from memoryos import MemoryClient

# Initialize with any AI platform
memory = MemoryClient(
    agent_id="my-assistant",
    user_id="user-123",
    api_key="mos_xxx"
```

```
)

# Store a memory
memory.remember(
    content="User mentioned they're training for a marathon",
    type="semantic",
    confidence=0.95,
    ttl="1y"  # Auto-expire after 1 year
)

# Retrieve relevant memories
context = memory.recall(
    query="What are user's fitness goals?",
    limit=5,
    recency_weight=0.3
)

# Forget on request (GDPR compliance)
memory.forget(user_id="user-123", scope="all")
```

Features: - Sub-10ms retrieval latency globally - Automatic memory consolidation (compress old memories) - Intelligent relevance ranking - Built-in embedding generation - Multi-modal support (text, images, audio transcripts)

2. MemoryOS Privacy Suite

Enterprise-grade privacy controls:

- **Consent Management** — Track what users consented to remember
- **Retention Policies** — Auto-delete after X days/months
- **Right to Delete** — One API call removes all user data
- **Data Residency** — Store memories in specific regions (EU, US, etc.)
- **Audit Logs** — Complete trail of memory access
- **PII Detection** — Auto-redact sensitive information

```
# Configure privacy for GDPR compliance
memory.configure_privacy(
    region="eu-west-1",
    retention_days=365,
    pii_redaction=True,
    consent_required=True
)
```

3. MemoryOS Sharing (Multi-Agent)

Enable agent coordination through shared memory:

- **Team Memory** — Multiple agents share context (support team, sales team)
- **Organization Memory** — Company-wide knowledge accessible to all agents
- **Cross-Platform Sync** — User's memory works across different AI providers
- **Memory Permissions** — Fine-grained access control

```
# Share memory across agent team
memory.share(
    memory_id="mem_xxx",
    with_agents=["support-agent-1", "support-agent-2"],
    permissions="read"
```

```
)

# Create organization memory
org_memory = MemoryClient(
    scope="organization",
    org_id="acme-corp"
)
org_memory.remember("Company holiday policy: ...", type="semantic")
```

4. MemoryOS Analytics

Understand how memory impacts agent performance:

- **Memory Utilization** — Which memories are accessed most?
- **Recall Quality** — Are retrieved memories relevant?
- **Coverage Gaps** — What questions can't be answered from memory?
- **User Insights** — Aggregate patterns (anonymized) across user base

---

Business Model

Pricing Tiers

Tier	Price	Memories/Month	Features
Free	$0	10K	Core API, 30-day retention
Pro	$49/mo	500K	Privacy suite, 1-year retention
Team	$199/mo	5M	Multi-agent sharing, analytics
Enterprise	Custom	Unlimited	Dedicated infra, SLA, on-prem option

Revenue Streams

1. **API Usage** — Per-memory storage and retrieval
2. **Platform Fees** — Monthly subscriptions
3. **Enterprise Contracts** — Annual deals with large deployments
4. **Memory Marketplace** — Commission on shared memory templates

Unit Economics

- **Cost to store 1M memories:** ~$5 (vector DB + blob storage)
- **Revenue per 1M memories:** ~$50 (Pro tier)
- **Gross margin:** 90%+

Market Analysis

TAM / SAM / SOM

Total Addressable Market (TAM): $30B by 2030 - All AI applications requiring persistent state - Includes chatbots, agents, copilots, autonomous systems

Serviceable Addressable Market (SAM): $8B by 2028 - Conversational AI and agent platforms specifically - ~100M deployed agents × $80/year average

Serviceable Obtainable Market (SOM): $400M by 2028 - 5% market share with first-mover advantage - Focus on enterprise and developer platforms

Competitive Landscape

Competitor	Weakness
OpenAI Memory	Platform-locked, basic, no privacy controls
Pinecone/Weaviate	Raw vector DBs, no memory semantics
LangChain Memory	Framework-specific, not standalone infra
Zep	Early, limited scale, narrow focus
Custom Solutions	Expensive, no standards, fragmented

Our moat: We're not a vector database — we're a cognitive memory layer with privacy-first design and multi-agent capabilities.

Go-To-Market Strategy

Phase 1: Developer Adoption (Months 1-6)

1. **Open Source Core** — Release basic memory primitives as OSS
2. **Developer Experience** — Best-in-class SDKs (Python, TypeScript, Go)
3. **Integration Partnerships** — LangChain, LlamaIndex, AutoGen
4. **Content Marketing** — "How to give your AI agent perfect memory"

Phase 2: Platform Expansion (Months 6-12)

1. **Framework Integrations** — Native plugins for major agent frameworks
2. **Cloud Marketplace** — AWS, Azure, GCP listings
3. **Enterprise Pilots** — 10 design partners in fintech, healthcare, support

Phase 3: Market Leadership (Months 12-24)

1. **Memory Standard** — Propose open memory format (like OpenAPI for memory)
2. **Enterprise Sales** — Dedicated team for $100K+ deals
3. **Platform Ecosystem** — Third-party memory providers and tools

---

Why Now?

Timing Catalysts

1. **Agent Explosion** — 2025-2026 is the year of AI agents. Thousands of companies building agents, all need memory.
2. **Context Limits Hit** — Even 200K tokens isn't enough. Memory systems are now essential, not optional.
3. **Enterprise AI Adoption** — Companies deploying AI at scale need persistent, compliant memory.
4. **Privacy Regulations** — GDPR, CCPA, AI Act all require data controls that memory systems must enforce.

5. **Multi-Agent Systems** — The rise of agent teams requires shared memory infrastructure.

Why This Team?

(Founding team to be assembled)

Ideal backgrounds: - Database infrastructure (Redis, Postgres, MongoDB veterans) - AI/ML platforms (former OpenAI, Anthropic, Google AI) - Enterprise sales (Snowflake, Datadog, Twilio alumni) - Privacy engineering (OneTrust, BigID experience)

Traction & Milestones

6-Month Goals

- ☐ Launch open-source memory SDK
- ☐ 1,000 developers using free tier
- ☐ 10 enterprise design partners
- ☐ $500K ARR

12-Month Goals

- ☐ 10,000 developers, 100 paying customers
- ☐ Launch privacy suite and sharing features
- ☐ Partnership with 2 major agent platforms
- ☐ $3M ARR

24-Month Goals

- ☐ 100K developers, 1,000 customers
- ☐ Memory standard proposal published
- ☐ Enterprise sales team (10 AEs)
- ☐ $20M ARR, path to profitability

Funding Ask

Seed Round: $5M

Use of Funds: - 40% Engineering (8 engineers) - 25% Infrastructure (global deployment) - 20% Go-to-market (DevRel, content, partnerships) - 15% Operations

Milestones to Series A: - $3M ARR - 10,000 developers - 50 enterprise customers - Clear path to memory standard adoption

---

Key Metrics

Metric	Target (12mo)
Developers	10,000
Memories Stored	10B
API Calls/Day	100M
Enterprise Customers	50
ARR	$3M

Metric	Target (12mo)
Net Revenue Retention	150%

Vision: The Memory Layer for AI

MemoryOS isn't just a product — it's the foundational infrastructure for AI that actually knows you.

Imagine: - Your AI assistant remembers every conversation, preference, and context — forever - Enterprise agents that learn from every customer interaction - Multi-agent teams that share knowledge seamlessly - Privacy controls that put users in charge of their data - A universal memory format that works across any AI platform

The future we're building: AI that doesn't just process — AI that remembers, learns, and grows.

Contact

Ready to give your AI agents perfect memory?

- Website: memoryos.ai
- Email: founders@memoryos.ai
- Twitter: @MemoryOS

"The best time to build AI memory infrastructure was yesterday. The second best time is now."

MemoryOS — Because AI Shouldn't Forget

Generated by The Godfather — February 9, 2026