

SimWorld AI — Synthetic Reality Generation for Physical AI

The “AWS” for AI Training Environments — \$50B+ Market Opportunity

Executive Summary

SimWorld AI is the foundational infrastructure for generating photorealistic, physics-accurate synthetic environments on demand. As humanoid robots, autonomous vehicles, drones, and embodied AI systems proliferate, they all face the same bottleneck: **training data from the physical world is expensive, dangerous, slow, and limited.**

SimWorld solves this by generating infinite variations of any real-world scenario — warehouses, city streets, factories, homes, construction sites — with perfect ground truth labels, controllable physics, and procedural variation.

The thesis: Every company building physical AI will need SimWorld. We become the picks-and-shovels provider for the robotics revolution.

The Problem

The Physical AI Data Crisis

The 2025-2026 robotics boom has created an unprecedented challenge:

1. **Real-world data is expensive** — Collecting one hour of warehouse robot training data costs \$10,000+ (equipment, facilities, human operators, labeling)
2. **Real-world data is dangerous** — Testing autonomous systems in real environments risks property damage, injury, and regulatory issues
3. **Real-world data is limited** — You can't easily create edge cases (a forklift crashing, a fire breaking out, unusual object configurations)
4. **Real-world data lacks ground truth** — Perfect segmentation, depth, and physics labels require expensive manual annotation
5. **Real-world conditions vary** — Lighting, weather, wear-and-tear, object arrangements change constantly. Models trained on limited real data fail in novel conditions

The Current State

- Tesla spends \$1B+/year on data collection and simulation for FSD
- Waymo has driven 20M+ real miles but still struggles with edge cases
- Figure and Boston Dynamics need vast training data for humanoid manipulation
- Amazon runs the largest fleet of warehouse robots but data collection is a massive cost center

Everyone is building their own simulation infrastructure. It's wasteful, expensive, and siloed.

The Solution

SimWorld AI Platform

Generate any physical environment, on demand, at infinite scale.

Core Capabilities 1. Procedural World Generation

```
"Generate 10,000 variations of a 50,000 sq ft warehouse with:  
- Random rack configurations (pallet racks, shelving, bins)  
- Varying lighting (fluorescent, natural skylights, emergency)  
- Different floor conditions (clean, wet, debris)  
- Random inventory levels and product types  
- 5-50 workers and forklifts moving realistically"
```

→ SimWorld generates all 10,000 environments in under 1 hour

2. Photorealistic Rendering - Neural radiance field (NeRF) + diffusion model hybrid rendering - Real-time ray tracing with physically-based materials - Sub-millimeter geometric accuracy - Sensor-accurate noise models (camera, LiDAR, radar, IMU)

3. Physics Simulation - Rigid body and soft body dynamics - Fluid simulation (water, dust, smoke) - Deformable objects (cardboard boxes, fabric, food) - Accurate friction, mass, and material properties

4. Perfect Ground Truth - Pixel-perfect semantic segmentation - Dense depth maps - Surface normals and material properties - 6DoF object poses - Full physics state at every timestep

5. Controllable Scenarios - Define edge cases programmatically - Inject failures, anomalies, and rare events - A/B test policy changes in simulation - Adversarial scenario generation

API-First Design

```
import simworld

# Create a warehouse environment
env = simworld.create(
    template="warehouse",
    size_sqft=50000,
    inventory_density=0.7,
    lighting="mixed",
    workers=15,
    forklifts=3
)

# Generate training data
for episode in range(10000):
    # Randomize conditions
    env.randomize(
        lighting_variance=0.3,
        object_placement_variance=0.5,
        weather=["clear", "humid", "dusty"]
    )

    # Run simulation with robot agent
    obs = env.reset()
    while not done:
        action = robot_policy(obs)
        obs, reward, done, info = env.step(action)

        # Perfect labels included
        rgb = obs["camera"]
        depth = obs["depth"]
        segmentation = obs["segmentation"]
```

```

lidar = obs["lidar_points"]

# Export for training
dataset = env.export(format="tfrecord")

```

Market Opportunity

Total Addressable Market: \$50B+ by 2030

Physical AI Training Infrastructure

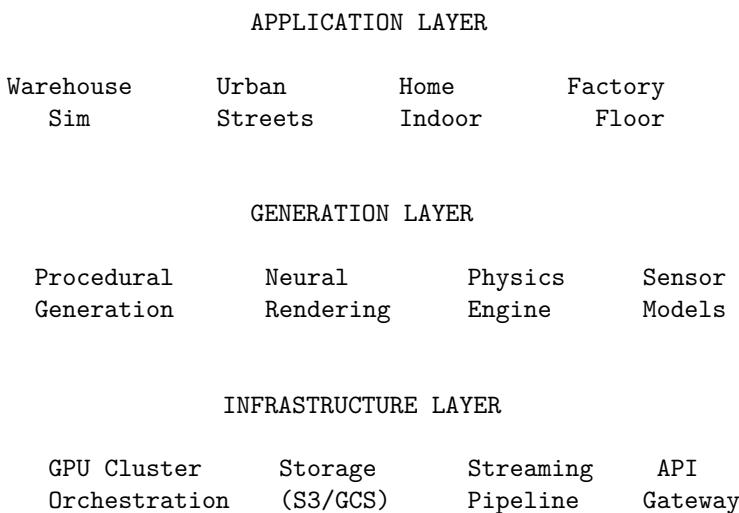
Segment	2026	2030	CAGR
Autonomous Vehicles	\$8B	\$25B	33%
Warehouse Robotics	\$2B	\$8B	41%
Humanoid Robots	\$500M	\$6B	85%
Industrial Automation	\$3B	\$10B	35%
Drones & UAVs	\$1B	\$4B	41%
Total	\$14.5B	\$53B	38%

Why Now?

- Generative AI breakthrough** — Diffusion models + NeRFs enable photorealistic generation at scale
 - Robotics inflection point** — Figure, Tesla Bot, 1X, Sanctuary shipping humanoids
 - Compute cost collapse** — GPU costs falling 30%/year makes large-scale simulation viable
 - Regulatory pressure** — AV companies required to prove safety in simulation before road testing
 - Consolidation opportunity** — Everyone building in-house, market ready for standardization
-

Product Architecture

Three-Layer Platform



Key Technical Innovations

- 1. Hybrid Neural-Procedural Generation** - Procedural rules for physical accuracy (no floating objects, correct physics) - Neural networks for photorealistic texturing and lighting - 100x faster than pure neural generation, 100x more realistic than pure procedural
 - 2. Domain Randomization Engine** - Automatically vary all visual and physical parameters - Curriculum learning support (easy → hard scenarios) - Distribution matching to real-world data statistics
 - 3. Sim-to-Real Transfer Module** - Style transfer to match target deployment environment - Uncertainty quantification for simulation confidence - Automatic domain gap detection and correction
 - 4. Scenario Mining** - Analyze real-world failures, generate similar scenarios - Adversarial scenario generation (attack the policy) - Coverage metrics for environment exploration
-

Business Model

Pricing Tiers

Starter — \$5,000/month - 100 GPU-hours of generation - 5 environment templates - Standard sensor models - Community support

Pro — \$25,000/month - 1,000 GPU-hours of generation - All environment templates - Custom sensor calibration - Priority support - API access

Enterprise — Custom (\$100K-\$1M+/year) - Unlimited generation - Custom environment development - On-premise deployment option - Dedicated success team - SLA guarantees

Revenue Model

Stream	Year 1	Year 3	Year 5
Platform Subscriptions	\$2M	\$30M	\$150M
Custom Environment Dev	\$1M	\$15M	\$50M
Professional Services	\$500K	\$10M	\$40M
Marketplace (Assets)	\$0	\$5M	\$30M
Total ARR	\$3.5M	\$60M	\$270M

Go-to-Market Strategy

Phase 1: Warehouse Robotics (Months 1-12)

Why start here: - Fastest-growing segment (Amazon, Walmart, Ocado all scaling) - Controlled environment (easier to simulate accurately) - Clear ROI (reduce real-world testing costs by 80%) - Strong network effects (assets reusable across customers)

Target customers: - Amazon Robotics - Locus Robotics - 6 River Systems (Shopify) - Berkshire Grey - Covariant

Launch strategy: 1. Partner with 2-3 design partners for beta (reduced price for feedback) 2. Publish benchmark showing sim-to-real transfer accuracy 3. Release warehouse asset pack on ROS/Isaac Sim marketplaces 4. Sponsor ICRA and CoRL conferences

Phase 2: Autonomous Vehicles (Months 12-24)

Why: - Largest market segment - Regulatory tailwinds (NHTSA requiring simulation validation) - High willingness to pay

Target customers: - Waymo, Cruise, Aurora - Tesla (partnership unlikely but they're the benchmark) - Motional, Argo successors - Chinese AV companies (WeRide, Pony.ai, Momenta)

Phase 3: Humanoid Robots (Months 18-36)

Why: - Fastest growth rate (85% CAGR) - Most complex simulation needs - Highest moat potential

Target customers: - Figure AI - Tesla Optimus - 1X Technologies - Sanctuary AI - Boston Dynamics

Competitive Landscape

Current Solutions (All Inadequate)

Solution	Limitation
NVIDIA Isaac Sim	Powerful but complex, no generation, requires in-house assets
Unity/Unreal	Gaming engines, not optimized for robotics, no domain randomization
Gazebo/PyBullet	Open source, low fidelity, no neural rendering
In-house (Tesla, Waymo)	\$100M+ investment, siloed, not productized
Synthesis AI	People/faces only, not environments
Datagen	Indoor scenes only, limited physics

SimWorld AI Differentiation

1. **End-to-end generation** — Don't just render assets, generate entire worlds
2. **Physics-first** — Accurate dynamics, not just pretty pictures
3. **API-native** — Built for ML pipelines, not game developers
4. **Infinite variation** — Procedural + neural = unlimited scenarios
5. **Sim-to-real focus** — Transfer learning built into the platform

Defensibility

- **Data network effects** — More customers = more assets = better generation
- **Domain expertise moat** — Deep robotics + ML + graphics intersection
- **Integration stickiness** — Training pipelines built around SimWorld
- **Benchmark ownership** — Become the standard for simulation fidelity

Team Requirements

Founding Team (Needed)

CEO — Robotics industry veteran, enterprise sales experience - Ex-Waymo, Tesla, Amazon Robotics leadership - Track record of selling to technical buyers

CTO — World-class graphics + ML researcher - Publications in SIGGRAPH, CVPR, CoRL - Experience with neural rendering, procedural generation - Ex-NVIDIA, DeepMind, Google Research

VP Engineering — Simulation infrastructure expert - Built large-scale simulation systems - Ex-Waymo Simulation, Tesla Dojo, NVIDIA Isaac

Key Hires (First 18 Months)

Role	Count	Focus
ML Engineers (Rendering)	4	Neural rendering, diffusion models
ML Engineers (Physics)	3	Dynamics, contact, deformation
Graphics Engineers	4	Real-time rendering, GPU optimization
Robotics Engineers	3	Sensor models, sim-to-real
Platform Engineers	4	API, infrastructure, scale
Solutions Engineers	2	Customer integration
Total	20	

Traction & Milestones

Pre-Seed (Current)

- Founding team assembled
- Technical prototype (warehouse generation)
- 2-3 design partner LOIs

Seed (\$5M, Month 0-12)

- Launch warehouse simulation product
- 5 paying customers
- \$500K ARR
- Publish sim-to-real benchmark paper
- 15-person team

Series A (\$25M, Month 12-24)

- Expand to autonomous vehicles
- 25 enterprise customers
- \$5M ARR
- Launch self-serve platform
- 50-person team

Series B (\$75M, Month 24-36)

- Full physical AI coverage (humanoids, drones, industrial)
- \$30M ARR
- International expansion
- Asset marketplace launch
- 150-person team

Financial Projections

5-Year Model

Metric	Year 1	Year 2	Year 3	Year 4	Year 5
Customers	10	40	120	300	600
ARR	\$3.5M	\$18M	\$60M	\$150M	\$270M
Gross Margin	65%	70%	75%	78%	80%
Team Size	20	50	120	250	400
Burn Rate	\$8M	\$18M	\$35M	\$50M	\$60M
Cash Position	\$4M	\$20M	\$70M	\$120M	\$180M

Unit Economics (Target)

- **ACV (Average Contract Value):** \$200K (Enterprise), \$25K (Pro)
 - **CAC:** \$50K (Enterprise), \$5K (Pro)
 - **LTV/CAC:** 8x+
 - **Gross Margin:** 75%+ at scale
 - **Net Revenue Retention:** 140%+ (expansion from usage growth)
-

The Big Picture

Why This Becomes a \$10B+ Company

1. **Every robot needs training data** — We're infrastructure for a \$500B robotics industry
2. **Winner-take-most dynamics** — Network effects from shared assets and benchmarks
3. **Expands with the market** — As robots proliferate, simulation needs multiply
4. **Platform potential** — Marketplace for 3D assets, scenarios, pre-trained models
5. **M&A optionality** — Strategic to NVIDIA, Google, Amazon, Tesla

The Vision

SimWorld AI becomes the **foundational infrastructure layer for physical AI** — the way AWS became foundational for cloud computing.

Every robot, every autonomous vehicle, every drone, every industrial automation system will train in SimWorld before deploying in the real world.

We don't just simulate reality. We generate infinite realities so AI can master the physical world.

Risks & Mitigations

Risk	Mitigation
NVIDIA bundles simulation	Focus on generation, not rendering; partner with NVIDIA
Big tech builds in-house	Move fast, establish standards, make switching costly
Sim-to-real gap persists	Heavy R&D investment, academic partnerships, customer feedback loops
Compute costs don't fall	Optimize for efficiency, explore custom silicon
Slow robotics adoption	Diversify across AV, drones, industrial (less correlated)

Call to Action

For Investors: SimWorld AI is raising a \$5M seed round to build the foundational team and launch our warehouse simulation product. Target close: Q2 2026.

For Customers: Join our design partner program. Get early access to world-class simulation infrastructure at a fraction of the cost of building in-house.

For Talent: We're assembling a world-class team at the intersection of graphics, ML, and robotics. If you want to build the infrastructure for the physical AI revolution, let's talk.

SimWorld AI — Generate Reality. Train Intelligence.

Contact: founders@simworld.ai

Location: San Francisco, CA

Founded: 2026