

AgentAudit — The SOC 2 for AI Agents

The Vision

AgentAudit is the enterprise compliance and certification platform that audits, certifies, and continuously monitors AI agent behavior — becoming the trusted standard for AI agent deployment in regulated industries.

Every enterprise wants to deploy AI agents. None of them can prove to their board, auditors, regulators, or customers that those agents are safe, compliant, and behaving as intended. AgentAudit closes that trust gap.

The Problem

The AI Agent Deployment Paradox

- **\$150B+ enterprise AI market** by 2027 — growing 37% annually
- **87% of enterprises** say autonomous AI agents are “critical” to their 2027 strategy
- **Only 12%** have deployed production AI agents at scale
- **#1 blocker:** “We can’t prove to regulators and auditors that agents are safe”

The Compliance Nightmare

Challenge	Enterprise Reality	Cost
No audit trail	AI agents make thousands of decisions — no consistent logging standard	Failed audits, regulatory fines
Behavior drift	Agent behavior changes over time with model updates	Compliance violations discovered months later
Shadow agents	Employees deploy unapproved AI tools	Data breaches, IP leakage
Vendor black boxes	Can’t audit third-party AI agent behavior	Liability when agents fail
Regulatory uncertainty	EU AI Act, state laws, industry regulations all diverging	Paralysis, missed opportunities

The Stakes Are Massive

- **Financial services:** \$1.7B in AI-related fines in 2025 alone
- **Healthcare:** HIPAA violations now include AI agent behavior
- **EU AI Act:** 6% of global revenue penalties for non-compliance
- **SOX auditors:** Now requiring AI system documentation
- **Cyber insurance:** Policies excluding “autonomous AI actions”

Why Current Solutions Fail

- **MLOps platforms:** Track model metrics, not agent behavior
- **GRC tools:** Designed for humans, not autonomous systems
- **AI security tools:** Focus on attacks, not compliance
- **Internal audit teams:** Don't have AI expertise, drowning in manual work
- **Consulting firms:** \$500/hour, no continuous monitoring

There is no SOC 2 for AI agents. No ISO standard. No trusted third-party certification that enterprises can show to regulators, customers, and boards.

The Solution

AgentAudit: Continuous Compliance for Autonomous AI

Core Platform:

AGENTAUDIT PLATFORM

DISCOVER	MONITOR	CERTIFY
• Agent inventory • Shadow AI detection • Risk scoring	• Behavior logging • Drift detection • Anomaly alerts	• Automated audits • Compliance reports • Trust badges

COMPLIANCE FRAMEWORKS

EU AI Act HIPAA SOC 2 NIST PCI GDPR FedRAMP

CERTIFICATION ENGINE

AgentAudit Certified Level 1/2/3 Industry Badges

Product Suite

1. AgentAudit Discovery Find every AI agent in your organization - Network-level agent detection (API patterns, model calls) - SaaS integration scanning (which apps use AI?) - Employee-deployed AI inventory - Risk scoring and prioritization - Shadow AI remediation workflows

2. AgentAudit Monitor Continuous behavioral monitoring for production agents - Universal logging SDK (works with any framework) - Decision audit trails with full context - Behavioral baseline establishment - Drift detection and alerting - Cost and performance tracking - Incident replay and root cause analysis

3. AgentAudit Certify Automated compliance certification - Pre-built compliance frameworks (EU AI Act, HIPAA, SOC 2) - Automated evidence collection - Gap analysis and remediation guidance - Continuous certification status - Auditor-ready reports - Public trust badges for customer confidence

4. AgentAudit Trust Registry The public registry of certified AI agents - Searchable database of certified agents - Certification history and status - Customer reviews and incident reports - Vendor comparison tools - API for automated vendor risk assessment

The AgentAudit Standard™

We're not just building software — we're creating the industry standard:

Certification Levels:

Level	Requirements	Use Case
Level 1: Basic	Logging, basic controls, incident response	Internal tools, low-risk
Level 2: Enhanced	Behavioral monitoring, drift detection, audit trails	Customer-facing, moderate risk
Level 3: Regulated	Full compliance mapping, continuous monitoring, third-party audit	Healthcare, finance, government
Level 4: Critical	Human oversight, explainability, real-time controls	Life safety, autonomous decisions

Industry Badges: - **HealthAI Certified** — HIPAA-compliant agent operations - **FinanceAI Certified** — SOX/FINRA-compliant autonomous decisions - **EU AI Act Compliant** — Full European regulatory alignment - **Enterprise Ready** — SOC 2 Type II equivalent for AI agents

Why Now?

1. Regulatory Tsunami (2026-2027)

- **EU AI Act** full enforcement begins August 2026
- **US state AI laws** (Colorado, California, Texas) taking effect
- **SEC guidance** on AI disclosure imminent
- **Industry regulators** (OCC, FDA, FTC) issuing AI-specific rules

2. Enterprise AI Agent Explosion

- **GPT-4o, Claude 3.5, Gemini 2** — agents are finally capable
- Every enterprise SaaS adding “AI agents”

- Internal agent development accelerating
- Multi-agent systems emerging in production

3. The Trust Gap

- **Boards** asking “how do we know the AI is safe?”
- **Auditors** requesting AI system documentation
- **Customers** demanding AI transparency
- **Insurers** excluding autonomous AI from policies

4. No Incumbent Solution

- GRC tools (OneTrust, ServiceNow) designed for human processes
- MLOps tools (MLflow, Weights & Biases) focused on models, not agents
- Security tools (Protect AI, Robust Intelligence) focused on attacks
- **Nobody owns “AI agent compliance”**

Business Model

Revenue Streams

1. Platform Subscription (Core)

Tier	Price	Agents	Features
Startup	\$2,500/mo	Up to 10	Discovery, basic monitoring
Business	\$15,000/mo	Up to 100	Full monitoring, 1 framework
Enterprise	\$75,000/mo	Unlimited	All features, custom frameworks
Regulated	Custom	Unlimited	Dedicated support, on-prem option

2. Certification Fees

- **Initial certification:** \$5,000 - \$50,000 (by complexity)
- **Annual renewal:** 50% of initial
- **Expedited certification:** 2x premium
- **Industry-specific badges:** Additional \$10,000-25,000

3. Auditor Marketplace

- Train and certify third-party auditors
- 15% platform fee on auditor engagements
- Premium listing fees for auditor firms

4. Trust Registry

- **Vendor listing:** Free (basic) / \$5,000/year (premium)
- **API access:** \$1,000-10,000/month for procurement teams

- **Lead generation:** Revenue share on enterprise deals

Unit Economics

- **CAC:** \$25,000 (enterprise sales)
 - **ACV:** \$180,000 (blended)
 - **LTV:** \$900,000+ (5-year retention)
 - **Gross margin:** 85%
 - **Payback period:** 4 months
-

Market Size

Total Addressable Market (TAM)

\$45B+ by 2030 - Enterprise GRC software: \$15B - IT audit services: \$12B - AI compliance (emerging): \$18B

Serviceable Addressable Market (SAM)

\$12B by 2028 - Enterprises deploying AI agents: 50,000+ - Average compliance spend per enterprise: \$200K+ - Certification and audit services: \$4B

Serviceable Obtainable Market (SOM)

\$500M by 2028 (Year 3) - 2,000 enterprise customers - 15% market share in early-mover segment

Go-to-Market Strategy

Phase 1: Establish the Standard (Months 1-12)

Target: Design partners who need compliance yesterday - Financial services firms under regulatory pressure - Healthcare organizations deploying AI - Government contractors with AI initiatives

Strategy: 1. **Free discovery tool** — find shadow AI, generate reports 2. **Design partnership program** — 20 enterprises shape the product 3. **Publish the AgentAudit Framework** — open specification 4. **Regulatory engagement** — brief EU, SEC, state regulators 5. **Big Four partnerships** — train auditors on our framework

Milestone: 50 paying customers, framework published

Phase 2: Platform Growth (Months 12-24)

Target: Fast followers seeing competitors get certified

Strategy: 1. **Trust Registry launch** — network effects kick in 2. **Vertical playbooks** — healthcare, finance, government 3. **Channel partnerships** — resellers, consultants, MSPs 4. **Self-serve tier** — startups building AI products 5. **Conference blitz** — RSA, Gartner, industry events

Milestone: 500 customers, recognized standard

Phase 3: Market Leadership (Months 24-36)

Target: Mainstream enterprise adoption

Strategy: 1. Auditor ecosystem — 1,000+ certified auditors 2. Procurement integration — become default vendor check 3. Insurance partnerships — certification = lower premiums 4. International expansion — EU-first, then APAC 5. Acquisition of competitors — consolidate the space

Milestone: 2,000+ customers, \$500M ARR

Competitive Landscape

Current Alternatives


Competitor	Approach	Limitation
OneTrust	GRC platform	Built for human processes, bolt-on AI
ServiceNow	IT workflow	No agent-specific monitoring
Weights & Biases	ML experiment tracking	Model-focused, not agent compliance
Protect AI	AI security	Attack surface, not compliance
Big Four	Manual audits	\$500/hour, not continuous
Internal teams	Custom solutions	Expertise gap, not scalable

Our Moats

- 1. Standard ownership — we define what “AI agent compliance” means
- 2. Network effects — Trust Registry creates switching costs
- 3. Data flywheel — behavioral patterns improve detection
- 4. Auditor ecosystem — trained professionals locked to our framework
- 5. Regulatory relationships — first to brief regulators, shape rules

Technology Architecture

Core Platform



Behavior Modeling Drift Detection Anomaly Analysis
Compliance Mapping Risk Scoring Pattern Recognition

COMPLIANCE LAYER

Framework Engine Evidence Collection Gap Analysis
Report Generator Certification Workflow Badge Issuance

TRUST REGISTRY

Public Ledger Certification Status Incident History
Vendor Profiles API Gateway Verification Service

Key Technical Capabilities

- 1. Universal Agent Logging** - Framework-agnostic SDK (LangChain, AutoGen, CrewAI, custom) - Structured decision logging with context preservation - Minimal performance overhead (<5ms per decision) - Privacy-preserving options (on-prem, tokenization)
- 2. Behavioral Analysis** - Baseline establishment from historical behavior - Statistical drift detection (distribution shift) - Semantic drift detection (LLM-powered analysis) - Anomaly scoring with explainability
- 3. Compliance Automation** - Machine-readable compliance frameworks - Automatic evidence mapping - Gap identification with remediation playbooks - Continuous certification status computation
- 4. Trust Infrastructure** - Tamper-proof audit logs (append-only, hashed) - Public verification of certifications - Revocation and incident disclosure - API for automated vendor assessment

Team Requirements

Founding Team (4-6 people)

CEO/Co-founder - Enterprise software sales background - Regulatory/compliance domain expertise - Network in financial services or healthcare

CTO/Co-founder - Distributed systems at scale - Security and compliance engineering - Experience with audit/logging systems

Head of Compliance - Big Four audit background - GRC software experience - Regulatory relationships

Head of Product - Enterprise B2B product experience - AI/ML product management - User research expertise

Early Hires (Months 1-12)

- 3-4 engineers (backend, platform)
- 2 compliance analysts
- 2 enterprise sales
- 1 solutions engineer
- 1 content/community lead

Scaling Team (Year 2)

- Engineering: 20+
 - Sales & CS: 15+
 - Compliance & Audit: 10+
 - Operations: 5+
-

Funding Strategy

Pre-Seed (\$2M)

Use of funds: - Founding team (4 people, 12 months) - MVP development - 10 design partners - Framework development

Target investors: Angels with enterprise/compliance experience

Seed (\$8M)

Use of funds: - Team to 25 - Product launch - 50 paying customers - Framework publication - Regulatory engagement

Target investors: Enterprise SaaS specialists (Unusual, Costanoa, Work-Bench)

Series A (\$30M)

Use of funds: - Team to 80 - Trust Registry launch - 500 customers - Auditor ecosystem - International expansion

Target investors: Tier 1 enterprise (Andreessen, Bessemer, Sequoia)

Key Risks & Mitigations

Risk	Mitigation
Regulation moves slowly	Focus on enterprise self-governance first; regulations will follow
Incumbents add features	Speed to standard ownership; network effects from registry
Enterprises build internally	Certification must be third-party to have value
AI agent adoption slows	Compliance applies to all AI systems, not just agents

Risk	Mitigation
Framework doesn't get adopted	Open-source framework, big Four partnerships

Success Metrics

Year 1

- 50 enterprise customers
- \$5M ARR
- Framework published and endorsed
- 20 auditors certified
- 3 regulatory briefings completed

Year 2

- 500 enterprise customers
- \$50M ARR
- Trust Registry live with 1,000+ agents
- 200 certified auditors
- Industry recognition as the standard

Year 3

- 2,000+ enterprise customers
 - \$500M ARR
 - IPO-ready metrics
 - Recognized legal/regulatory standard
 - International presence (EU, APAC)
-

The Vision: Trust Infrastructure for the AI Economy

In 5 years, no enterprise will deploy an AI agent without AgentAudit certification. Our badge will be as ubiquitous as “SOC 2 Certified” or “HIPAA Compliant.”

We’re not just building a compliance tool — we’re building the **trust infrastructure for autonomous AI**.

When regulators need to understand AI agent governance → AgentAudit framework
 When boards ask “is our AI safe?” → AgentAudit certification
 When customers evaluate AI vendors → AgentAudit Trust Registry
 When insurers price AI risk → AgentAudit risk scores

The AI economy needs a trust layer. We’re building it.

Why This Team, Why Now

Timing is Perfect

- EU AI Act enforcement in 3 months
- Enterprise AI agent adoption inflecting
- No incumbent has claimed this space
- Regulatory pressure creating urgency

Unfair Advantages

- First to publish comprehensive framework
- Regulatory relationships from Day 1
- Big Four partnership strategy
- Network effects from Trust Registry

Exit Potential

- **IPO:** \$10B+ market cap as category leader
- **Strategic acquisition:** ServiceNow, Salesforce, Microsoft
- **Private equity:** Platform consolidation play

The AI agent era is here. Trust isn't optional — it's infrastructure.

AgentAudit: The SOC 2 for AI Agents.