

SelfCloud — Your AI Memory Layer for Every AI You Use

One Context. Every AI. Finally, AI That Actually Knows You.

The \$50B Problem: AI Context Fragmentation

Every knowledge worker in 2026 uses 5-10 different AI tools daily: - ChatGPT for brainstorming - Claude for writing - Copilot for coding - Gemini for research - Perplexity for search - Custom enterprise agents - Vertical-specific AI tools

The catastrophic problem? None of them talk to each other.

Every single AI conversation starts from zero. You repeat your preferences, your context, your projects, your communication style—**over and over and over again**.

The average knowledge worker wastes **47 minutes per day** re-establishing context with AI tools. That's **195 hours per year** of pure friction.

For enterprises, it's worse: institutional knowledge is scattered across thousands of disconnected AI conversations, creating massive liability, compliance gaps, and lost intellectual property.

The Insight: AI Needs a Personal Memory Layer

What if every AI you used could tap into a single, privacy-first memory layer that contained: - Your communication preferences - Your ongoing projects and goals - Your expertise and knowledge domains - Your organizational context - Your relationships and network - Your decision history and reasoning patterns

SelfCloud is the universal context protocol for personal and enterprise AI.

We're not building another AI. We're building the **memory infrastructure** that makes every AI dramatically better.

How SelfCloud Works

For Individuals

- 1. Connect Once, Remember Forever**
 - Install the SelfCloud layer (browser extension, API, or native app)
 - Every AI interaction automatically enriches your personal context graph
 - Consent-first: you control exactly what gets stored and shared
- 2. Universal Context Injection**
 - When you start a new AI conversation, SelfCloud injects relevant context
 - "Writing a product brief? Here's your company's tone guide, your last 3 briefs, and your product positioning."
 - Works with any AI—ChatGPT, Claude, Gemini, enterprise tools, everything
- 3. Smart Memory Management**
 - Auto-summarization of key decisions and insights
 - Forgetting by design: set retention policies
 - Export your entire AI memory anytime


For Enterprise

- 1. Institutional Memory Layer**
 - Capture collective knowledge from all AI interactions

- Maintain compliance and audit trails
 - Prevent knowledge loss from employee turnover
2. **Context Governance**
 - Control what context flows to which AI systems
 - Enforce data classification policies
 - Prevent sensitive information leakage
 3. **AI Asset Management**
 - Unified dashboard for all enterprise AI usage
 - Cost optimization across AI providers
 - Performance analytics and ROI measurement
-

The Technical Architecture

Privacy-First Design



Core Technical Components

1. **Context Capture Engine**
 - Browser extension intercepts AI conversations
 - API hooks for enterprise applications
 - Natural language processing for context extraction
 - Semantic understanding of user intent and preferences
2. **Privacy Fabric**
 - On-device processing by default

- Differential privacy for any cloud operations
 - User-controlled encryption keys
 - Granular consent management
3. **Context Graph**
- Knowledge graph of user's digital identity
 - Temporal context (what's relevant now vs. historically)
 - Relationship mapping (people, projects, organizations)
 - Preference learning (communication style, decision patterns)
4. **Universal Injection Protocol**
- OpenContext standard for AI context sharing
 - Provider-agnostic middleware
 - Graceful degradation for non-integrated AIs
 - Real-time context relevance scoring

Market Opportunity

Total Addressable Market: \$47B by 2030

Segment	Market Size	SelfCloud Opportunity
AI Productivity Tools	\$23B	Context layer = 15% premium
Enterprise Knowledge Management	\$12B	AI-native replacement
Personal Data Vaults	\$4B	AI-specific use case
AI Security & Governance	\$8B	Compliance infrastructure

Serviceable Addressable Market: \$12B

- 500M knowledge workers globally
- 30% will use multiple AI tools by 2028
- \$80/user/year average revenue potential

Beachhead Market: \$800M

- US enterprise early adopters (Fortune 2000)
- 5M seats at \$160/seat/year enterprise pricing

Business Model

Consumer: Freemium

Tier	Price	Features
Free	\$0	3 AI integrations, 1GB context storage, basic memory
Pro	\$12/month	Unlimited integrations, 50GB storage, advanced memory, export
Family	\$25/month	6 accounts, shared family context, parental controls

Enterprise: Seat-Based + Platform

Tier	Price	Features
Team	\$15/user/month	Up to 100 users, shared workspaces, basic governance
Business	\$30/user/month	Unlimited users, SSO, compliance, audit trails
Enterprise	Custom	On-prem option, custom integrations, dedicated support

Platform Revenue

- **API Access:** \$0.001 per context injection for third-party AI providers
- **Data Insights:** Anonymized, aggregated insights for AI providers (opt-in)
- **Partner Ecosystem:** Revenue share with integrated applications

Competitive Landscape

Current Solutions (All Inadequate)

Competitor	Approach	Why It Fails
Mem.ai	Personal knowledge base	Not AI-native, manual organization
Notion AI	Single-platform memory	Walled garden, limited context
Rewind.ai	Full capture	Privacy nightmare, no injection
ChatGPT Memory	Platform-specific	Only works with ChatGPT
Custom GPTs	Pre-configured personas	Static, doesn't learn

SelfCloud's Unfair Advantages

1. **Universal Protocol:** Works with ANY AI, not locked to one provider
2. **Privacy-First:** On-device processing, user-owned keys
3. **Enterprise-Ready:** Compliance, governance, audit from day one
4. **Network Effects:** More users = better context understanding
5. **First-Mover:** No one has built the universal context layer yet

Go-to-Market Strategy

Phase 1: Developer & Power User Adoption (Months 1-6)

Target: AI early adopters, developers, productivity enthusiasts

Tactics: - Launch on Product Hunt, Hacker News - Open-source the context capture SDK - Partner with AI newsletter creators (The Rundown, AI Breakfast) - Build integrations for top 10 AI tools - Free tier for developers

Milestone: 100K users, 10K paying subscribers

Phase 2: Enterprise Pilot (Months 6-12)

Target: Fortune 2000 innovation teams, consulting firms, tech companies

Tactics: - Direct sales to CIOs and Chief AI Officers - SOC 2 Type II certification - Pilot programs with 3-5 lighthouse customers - Case study development - Enterprise partner channel (Accenture, Deloitte)

Milestone: 20 enterprise customers, \$2M ARR

Phase 3: Platform Expansion (Months 12-24)

Target: Mass market, global enterprise

Tactics: - Launch OpenContext protocol as open standard - Partner integrations with major AI providers
- International expansion (EU, APAC) - Vertical-specific solutions (legal, healthcare, finance) - M&A of complementary tools

Milestone: 1M users, 200 enterprise customers, \$20M ARR

Financial Projections

Revenue Forecast

Year	Users	Enterprise Seats	ARR
2026	100K	5K	\$2M
2027	500K	50K	\$15M
2028	2M	200K	\$60M
2029	8M	800K	\$200M
2030	25M	2M	\$500M

Unit Economics

- **Consumer CAC:** \$15 (organic + content marketing)
- **Consumer LTV:** \$144 (12 months average, 35% convert to paid)
- **LTV:CAC Ratio:** 9.6x
- **Enterprise CAC:** \$25K (inside sales + pilots)
- **Enterprise LTV:** \$180K (3-year average contract, 95% retention)
- **LTV:CAC Ratio:** 7.2x

Funding Requirements

Round	Amount	Use of Funds	Milestone
Seed	\$4M	Core product, founding team	50K users, product-market fit
Series A	\$20M	GTM, enterprise, integrations	\$5M ARR, 100 enterprise customers
Series B	\$60M	Scale, international, platform	\$50M ARR, category leadership

Team Requirements

Founding Team (To Recruit)

Role	Profile
CEO	Enterprise SaaS founder, AI/ML background, 2nd-time founder
CTO	Ex-Google/Anthropic, privacy engineering, distributed systems
CPO	Consumer product leader, knowledge management or productivity
Head of AI	Research scientist, NLP/knowledge graphs, PhD preferred

Early Hires (First 20)

- 8 Engineers (infra, ML, frontend, mobile)
- 3 Product/Design
- 4 GTM (sales, marketing, partnerships)
- 2 Operations (finance, people)
- 3 Customer Success

Risk Analysis

Technical Risks

Risk	Likelihood	Impact	Mitigation
AI providers block integration	Medium	High	Open standard, user demand, legal strategy
Privacy breach	Low	Critical	On-device first, security audit, insurance
Context injection latency	Medium	Medium	Edge computing, caching, optimization

Market Risks

Risk	Likelihood	Impact	Mitigation
AI providers build native memory	High	High	Universal protocol, multi-provider value
Enterprise AI consolidation	Medium	Medium	Provider-agnostic positioning
Consumer privacy concerns	Medium	Medium	Transparency, user control, certification

Competitive Risks

Risk	Likelihood	Impact	Mitigation
Well-funded copycat	High	Medium	First-mover, network effects, enterprise moat
Big tech acquisition of competitor	Medium	Medium	Become acquisition target, build defensibility

Why Now?

The Convergence Moment

1. **Multi-AI Reality:** 2026 is the year everyone uses multiple AI tools daily
2. **AI Agent Explosion:** Autonomous agents need persistent context to function
3. **Enterprise AI Maturity:** Companies moving from experimentation to deployment
4. **Privacy Regulation:** GDPR, state laws create demand for user-controlled data
5. **Infrastructure Gap:** The picks-and-shovels play for AI is still wide open

The Window

Major AI providers are focused on making their individual products better. They have **zero incentive** to build interoperability.

By the time they realize the context layer is the most valuable piece, SelfCloud will own the standard.

The Vision: Own Your AI Relationship

Today, your relationship with AI is fragmented, ephemeral, and owned by corporations.

SelfCloud creates a future where:

- Your AI context belongs to YOU, forever
- Every AI you interact with knows you as well as your best colleague
- Switching AI providers is seamless—your memory comes with you
- Your accumulated AI interactions become a valuable personal asset
- Enterprises maintain institutional memory even as AI systems evolve

We're not building an AI company. We're building the identity layer for the AI age.

The Ask

Raising \$4M Seed Round

- Lead: \$2.5M from AI-focused fund
- Strategic: \$1M from AI provider or enterprise software leader
- Angels: \$500K from AI founders and operators

Use of Funds: - 50% Engineering (core product) - 25% GTM (launch and early traction) - 15% Operations - 10% Reserve

Target Close: Q3 2026

Appendix

A. Technical Specifications

Context Graph Schema:

```
{
  "user_id": "encrypted_uuid",
  "context_nodes": [
    {
      "type": "preference",
      "domain": "communication_style",
      "value": "concise, direct, data-driven",
      "confidence": 0.94,
      "last_updated": "2026-05-14T12:00:00Z",
      "source_count": 147
    },
    {
      "type": "project",
      "name": "Q2 Product Launch",
      "status": "active",
      "related_entities": ["marketing_team", "product_v2", "launch_date_june15"],
      "context_relevance": 0.87
    }
  ],
  "relationships": [
    {
      "from": "user",
      "to": "entity:sarah_chen",
      "type": "colleague",
      "context": "direct_report, marketing_lead"
    }
  ]
}
```

API Example:

```
from selfcloud import SelfCloudClient

# Initialize with user's encrypted key
client = SelfCloudClient(user_key="...")

# Get relevant context for a new AI conversation
context = client.get_context(
    prompt="Help me write a product announcement email",
    ai_provider="claude",
    max_tokens=2000
)

# Context includes: communication style, product details, past announcements,
# team members to CC, company tone guidelines
```

B. Privacy Architecture

Data Flow: 1. AI conversation captured locally 2. On-device NLP extracts context elements 3. Sensitive info filtered based on user rules 4. Context elements encrypted with user key 5. Encrypted sync to SelfCloud (optional) 6. Context injection uses relevance scoring 7. No raw conversation data ever leaves device

Compliance: - GDPR Article 20 (data portability): Full export - CCPA: Do not sell, user control - SOC 2 Type II: In progress - HIPAA: Enterprise tier available

C. Integration Roadmap

Launch (Q3 2026): - ChatGPT (via browser extension) - Claude (via browser extension) - Gemini (via browser extension) - Perplexity (via browser extension)

Q4 2026: - Copilot (VS Code extension) - Notion AI (API integration) - Custom enterprise agents (SDK)

2027: - Native mobile (iOS, Android) - Slack AI, Teams AI - Salesforce Einstein - OpenContext protocol v1.0

“The AI that knows you best will win. SelfCloud makes sure that AI is every AI you choose.”

Contact: founders@selfcloud.ai

Website: selfcloud.ai (launching Q3 2026)