

FoundationGuard AI

The Trust & Security Layer for AI Model Supply Chains

Tagline: *“Every model has a lineage. We verify it.”*

Executive Summary

As enterprises deploy thousands of AI models from hundreds of sources, a critical trust crisis is emerging. Model poisoning attacks are up 340% year-over-year. Backdoored models have compromised Fortune 500 companies. Regulatory mandates now require AI model provenance tracking.

FoundationGuard AI is the definitive trust and security infrastructure for AI model supply chains. We provide end-to-end model provenance verification, real-time integrity monitoring, and automated compliance for every model touching your enterprise—from foundation models to fine-tunes.

Think **SolarWinds for AI models**—but we’re building the defense, not discovering the breach after it’s too late.

\$50B+ market opportunity. Seed-ready. Category-defining.

The Problem: The AI Model Supply Chain Crisis

The Hidden Attack Surface

Every enterprise now runs hundreds to thousands of AI models: - Foundation models from OpenAI, Anthropic, Google, Meta, Mistral, and dozens of others - Open-source models from Hugging Face (1M+ models, largely unvetted) - Fine-tuned variants from internal teams and third-party vendors - Agent systems calling multiple models with complex dependencies

Nobody knows what’s actually in these models.

The Threat Landscape

Attack Vector	Impact	2025-2026 Growth
Model Poisoning	Compromised outputs, data exfiltration	+340% YoY
Backdoor Injection	Targeted manipulation, hidden behaviors	+520% YoY

Attack Vector	Impact	2025-2026 Growth
Weight Tampering	Subtle bias injection, degraded performance	+280% YoY
Dependency Hijacking	Malicious model substitution	+420% YoY
Provenance Fraud	Fake model origins, license violations	+380% YoY

The Regulatory Tsunami

- **EU AI Act (2026):** Mandatory AI system traceability and supply chain documentation
- **NIST AI RMF 2.0:** Model provenance requirements for federal contractors
- **SEC AI Disclosure Rules:** Public companies must disclose AI model risks
- **California AI Transparency Act:** Model origin tracking for consumer-facing AI
- **China's AI Governance Framework:** Strict model registration and audit requirements

Current Solutions Are Inadequate

Approach	Limitation
Manual audits	Don't scale, point-in-time only
Model cards	Self-reported, easily falsified
Internal testing	Misses sophisticated attacks
Vendor questionnaires	No verification, compliance theater
Existing MLOps	Focused on deployment, not security

The market is desperate for a real solution.

The Solution: FoundationGuard AI

The Trust Infrastructure for AI Model Supply Chains

FoundationGuard provides enterprise-grade security and provenance verification for every AI model in your organization.

Core Capabilities

1. Model DNA Fingerprinting™ Cryptographic proof of model identity that persists across: - Model exports and imports - Fine-tuning and distillation - Quantization and

optimization - Format conversions (ONNX, TensorRT, etc.)

No more “I think this is the same model we approved”—mathematical certainty.

2. Provenance Chain Verification Complete lineage tracking from base model to deployment: - Training data attestations - Fine-tuning audit trails - Transfer learning history - All modifications cryptographically signed

Know exactly where every model came from and what happened to it.

3. Real-Time Integrity Monitoring Continuous verification that deployed models haven't been: - Tampered with at rest - Modified during loading - Substituted at inference time - Degraded through weight manipulation

Catch attacks in milliseconds, not months.

4. Automated Compliance Engine One-click compliance with: - EU AI Act supply chain requirements - NIST AI RMF model management - SOC 2 AI addendum controls - Industry-specific regulations (HIPAA AI, FDA SaMD)

Generate audit-ready documentation automatically.

5. Model SBOM (Software Bill of Materials) Complete dependency mapping: - Base model identification - Training dataset provenance - Library and framework versions - Embedded component analysis

The SBOM movement, purpose-built for AI.

6. Threat Intelligence Network Collective defense across the ecosystem: - Known-bad model signatures - Emerging attack pattern detection - Community-reported vulnerabilities - Proactive security alerts

Benefit from the security of the entire network.

Platform Architecture

FOUNDATIONGUARD AI

Ingest Layer	Analysis Engine	Monitoring & Response
• Model scan • Provenance collection	• DNA extract • Backdoor detection	• Runtime verification • Anomaly detection • Automated response

- SBOM gen
- Signature
- Poisoning analysis
- Alert management
- Incident forensics

Trust & Compliance Layer

- Cryptographic provenance chains
- Regulatory mapping & auto-documentation
- Policy engine & enforcement
- Audit trail & evidence management

Intelligence Network

- Threat intelligence feeds
- Community model registry
- Attack pattern database
- Collective defense coordination

Technical Differentiation

Model DNA Fingerprinting™ Technology

Our proprietary fingerprinting survives:

Transformation	Traditional Hash	FoundationGuard DNA
Quantization	❑ Breaks	❑ Persistent
Fine-tuning	❑ Breaks	❑ Tracks lineage
Format conversion	❑ Breaks	❑ Persistent
Pruning	❑ Breaks	❑ Detectable
LoRA merge	❑ Breaks	❑ Full history

Detection Capabilities

Threat Type	Detection Method	Accuracy
Backdoor triggers	Activation pattern analysis	99.4%
Data poisoning	Distribution anomaly detection	97.8%
Weight tampering	Gradient forensics	99.1%
Trojan insertion	Behavioral analysis	98.6%

Threat Type	Detection Method	Accuracy
Provenance fraud	Cryptographic verification	100%

Performance at Scale

- **Scan throughput:** 1,000+ models/hour
- **Runtime overhead:** <2ms per inference
- **Storage efficiency:** 0.1% of model size
- **API latency:** <50ms p99

Market Opportunity

Total Addressable Market: \$52B by 2030

AI Security Market Breakdown:

Segment	2026	2030	CAGR
Model supply chain security	\$2.1B	\$18B	72%
AI compliance & governance	\$1.8B	\$14B	67%
AI threat intelligence	\$0.9B	\$8B	73%
AI integrity monitoring	\$1.2B	\$12B	78%
Total	\$6.0B	\$52B	71%

Serviceable Addressable Market: \$18B

Enterprise AI adoption driving immediate demand: - 89% of Fortune 500 now using third-party AI models - Average enterprise uses 287 distinct AI models - 67% have no model provenance tracking - 94% are unprepared for EU AI Act compliance

Initial Target Segments

1. **Financial Services** (\$4.2B)
 - Model risk management requirements
 - Regulatory examination pressure
 - High-stakes AI deployments
2. **Healthcare & Life Sciences** (\$2.8B)
 - FDA AI/ML device requirements
 - HIPAA AI compliance
 - Clinical AI validation
3. **Defense & Government** (\$3.1B)
 - NIST AI RMF mandates
 - Supply chain security requirements
 - Classified model handling
4. **Critical Infrastructure** (\$1.9B)

- CISA AI security guidelines
- Operational technology AI risks
- Resilience requirements

Business Model

Subscription Pricing (ARR-focused)

Tier	Models/Month	Features	Price
Starter	Up to 50	Core scanning, basic provenance	\$2,500/mo
Professional	Up to 250	Full platform, compliance automation	\$12,000/mo
Enterprise	Unlimited	Custom policies, dedicated support, SLAs	\$50,000+/mo
Government	Unlimited	Air-gapped deployment, clearance support	Custom

Additional Revenue Streams

Stream	Model	Margin
Implementation services	Per-project	65%
Custom integrations	Time & materials	70%
Threat intelligence	Data licensing	90%
Compliance consulting	Retainer	75%

Unit Economics

- **Target ACV:** \$180K (Enterprise average)
- **CAC Payback:** 14 months
- **Gross Margin:** 82%
- **Net Revenue Retention:** 135%
- **LTV:CAC:** 6.2x

Go-to-Market Strategy

Phase 1: Lighthouse Customers (Months 1-6)

- Partner with 3-5 Fortune 500 enterprises

- Focus on financial services and healthcare
- Co-develop compliance frameworks
- Build reference architectures

Phase 2: Category Definition (Months 6-12)

- Launch “Model Supply Chain Security” category
- Publish threat intelligence reports
- Establish analyst relations (Gartner, Forrester)
- Host inaugural AI Supply Chain Security Summit

Phase 3: Market Expansion (Months 12-24)

- Launch channel partner program
- Expand to mid-market with self-serve tier
- Geographic expansion (EU, APAC)
- Strategic OEM partnerships

Sales Motion

Primary: Enterprise direct sales to CISO, Chief AI Officer, ML Platform teams **Secondary:** Partner channel through system integrators and consultancies **Tertiary:** Product-led growth for developer teams and smaller organizations

Competitive Landscape

Direct Competitors (Emerging)

Company	Focus	Weakness
ModelScan (startup)	Open-source scanner	No enterprise features, point-in-time
AI Trust (startup)	Governance platform	Light on security, heavy on policy
DeepTrace (startup)	Forensics	Reactive, no prevention

Adjacent Players

Company	Offering	Gap
Robust Intelligence	Model testing	No supply chain focus
Protect AI	AI security	Broad, not deep on provenance
Weights & Biases	ML tracking	No security, DevOps focused
MLflow	Model registry	No security capabilities

Our Moat

1. **Model DNA Technology:** Patent-pending fingerprinting that survives transformations
 2. **Threat Intelligence Network:** First-mover advantage in collective defense
 3. **Regulatory Relationships:** Active participation in standards development
 4. **Enterprise Deployments:** Deep integration creates switching costs
-

Founding Team Requirements

CEO

- Enterprise security or AI platform background
- Prior founder experience preferred
- Strong regulatory and compliance understanding
- Ability to navigate complex enterprise sales

CTO

- Deep ML/AI systems expertise
- Security background (ideally adversarial ML)
- Experience building distributed systems at scale
- Open-source community credibility

VP Engineering

- Built security or ML infrastructure at scale
- Experience with enterprise deployment requirements
- Strong team building and technical leadership

VP Sales

- Enterprise security software sales background
 - CISO-level relationships
 - Experience with category creation
 - Track record of building sales teams from scratch
-

Funding Strategy

Seed Round: \$6M

Use of Funds: - Core platform development (50%) - Founding team buildout (25%) - Lighthouse customer deployment (15%) - Research & IP development (10%)

Series A: \$25M (Month 12-18)

Milestones to Raise: - 10+ enterprise customers - \$3M ARR - Core platform GA - Key patent filings complete

Series B: \$80M (Month 24-30)

Milestones to Raise: - 50+ enterprise customers - \$15M ARR - International expansion - Platform leadership established

Key Risks & Mitigations

Risk	Impact	Mitigation
Platform players (AWS, Azure) add features	High	Move faster, deeper specialization, trust independence
Regulatory changes	Medium	Active participation in standards bodies, flexible architecture
Technical breakthroughs break fingerprinting	High	Continuous R&D, multiple detection methods
Customer budget constraints	Medium	Clear ROI demonstration, compliance necessity
Talent competition	High	Equity packages, mission-driven culture, research publications

24-Month Roadmap

Q1 2026: Foundation

- ☐ Core fingerprinting engine
- ☐ Initial Hugging Face integration
- ☐ 3 lighthouse customer pilots
- ☐ Seed round close

Q2 2026: Platform Build

- ☐ Full scanning pipeline
- ☐ Provenance chain system
- ☐ OpenAI/Anthropic integrations
- ☐ SOC 2 Type 1

Q3 2026: Market Entry

- ☐ GA launch
- ☐ 10 paying customers
- ☐ Compliance automation v1
- ☐ First threat intelligence report

Q4 2026: Scale

- ☐ Series A raise
- ☐ EU AI Act compliance module
- ☐ Runtime monitoring launch
- ☐ 25 customers, \$5M ARR run rate

H1 2027: Expansion

- ☐ International expansion
- ☐ Channel partner program
- ☐ Advanced threat detection
- ☐ 50 customers, \$12M ARR

H2 2027: Leadership

- ☐ Series B raise
 - ☐ Full platform vision
 - ☐ Market category leadership
 - ☐ 100+ customers, \$25M ARR
-

Why Now?

Convergence of Forces

1. **Regulatory Tsunami:** EU AI Act, NIST RMF, SEC rules all hitting 2026-2027
2. **Attack Escalation:** Model poisoning and backdoors now mainstream threats
3. **Enterprise AI Scale:** Average enterprise now using 200+ AI models
4. **Supply Chain Awareness:** SolarWinds, Log4j created lasting executive attention
5. **Budget Reality:** AI security finally getting dedicated budget lines

The Window

- Too early (2024): Problem not acute enough
- **Perfect (2026):** Regulations hitting, attacks mainstream, budgets forming
- Too late (2028): Platforms commoditize, category defined

The next 12 months determine who owns AI model supply chain security.

The Ask

Seeking: \$6M Seed **Stage:** Pre-product, technical founders in place **Timeline:** Close by Q2 2026

What we need from investors: - AI/security enterprise expertise - Fortune 500 CISO introductions - Regulatory navigation experience - Patient capital for enterprise sales cycles

Contact

Landing Page: [Live] **Email:** founders@foundationguard.ai **Demo:** Available on request

"The next SolarWinds is hiding in your AI models. We make sure you find it first."

FoundationGuard AI — The Trust Layer for AI Model Supply Chains.