

AgentCredit — The FICO Score for AI Agents

“In a world of 10 billion AI agents, who do you trust?”

Executive Summary

As autonomous AI agents proliferate across enterprise and consumer applications, a critical infrastructure gap emerges: **there’s no standardized way to evaluate agent trustworthiness**. AgentCredit creates the universal reputation and credit scoring system for AI agents — enabling trust-based commerce, risk management, and accountability in the agentic economy.

The Opportunity: By 2027, analysts project 10+ billion active AI agents handling transactions, negotiations, and decisions worth trillions of dollars. Yet there’s no FICO, no Dun & Bradstreet, no Moody’s for AI agents. AgentCredit fills this infrastructure gap.

The Problem

The Trust Crisis in Agentic AI

- 1. No Accountability** - AI agents operate across organizations with no persistent reputation - Bad actors can spin up new agents after failures - No historical track record for decision-making quality
- 2. Enterprise Risk Exposure** - Companies deploying agent-to-agent commerce have no way to vet counterparty agents - Insurance and liability gaps for autonomous transactions - Regulatory pressure building for AI accountability
- 3. Agent Fraud & Manipulation** - Compromised agents can impersonate legitimate ones - Prompt injection attacks can hijack agent behavior - No forensic capability to trace agent misconduct
- 4. Consumer Vulnerability** - Personal AI assistants interact with unknown commercial agents - No transparency into agent objectives or alignment - Privacy and manipulation risks unaddressed

Market Pain Points

Stakeholder	Problem	Cost
Enterprises	Can’t trust third-party AI agents	\$500B+ in delayed AI adoption
Insurers	Can’t price AI agent liability	Trillions in unpriced risk
Regulators	No visibility into agent behavior	Mounting political pressure
Consumers	Exploited by unaccountable agents	Trust erosion, backlash

The Solution

AgentCredit: Trust Infrastructure for the Agentic Economy

Core Platform Components:

1. Universal Agent Identity (UAI)

- Cryptographic identity binding for AI agents
- Immutable provenance chain (model, training, operator, purpose)
- Cross-platform identity federation
- Hardware attestation for edge agents

2. AgentCredit Score (ACS) A multi-dimensional trust score (300-850 scale, like FICO):

- **Reliability Score:** Task completion rate, latency, uptime
- **Accuracy Score:** Factual correctness, hallucination rate
- **Safety Score:** Boundary adherence, policy compliance
- **Ethics Score:** Fairness metrics, bias detection
- **Security Score:** Compromise history, vulnerability exposure

3. Behavioral Observatory

- Real-time monitoring across participating platforms
- Pattern detection for anomalous behavior
- Forensic capabilities for incident investigation
- Privacy-preserving computation (no raw data exposure)

4. Trust API & SDKs

```
// Check agent trustworthiness before interaction
const score = await AgentCredit.verify(agentId);
if (score.overall < 650 || score.safety < 600) {
  reject("Agent does not meet trust threshold");
}
```

```
// Report interaction outcome
await AgentCredit.report({
  agentId,
  interaction: "contract_negotiation",
  outcome: "successful",
  accuracy: 0.95,
  compliance: true
});
```

5. Agent Insurance Integration

- Pre-qualified agents for coverage
- Dynamic premium pricing based on ACS
- Automated claims processing for agent failures
- Liability allocation for multi-agent incidents

Business Model

Revenue Streams

Stream	Model	Year 3 Target
Score API	Per-query pricing (\$0.001-\$0.10)	\$80M ARR

Approach	Limitations
Internal Monitoring	Siloed, no cross-platform visibility
API Gateways	Security-focused, not trust/reputation
AI Auditing Firms	Point-in-time, expensive, not real-time
Blockchain Identity	Identity only, no behavioral scoring

AgentCredit Differentiation

1. **Universal Coverage:** Cross-platform, cross-vendor
2. **Real-Time Scoring:** Continuous, not periodic
3. **Behavioral Intelligence:** Actual performance, not capability claims
4. **Insurance Integration:** Financial accountability built-in
5. **Network Effects:** Value increases with participation

Competitive Moat

- **Data Network Effects:** More participants = better scores
- **Standard-Setting:** Push for regulatory adoption
- **Integration Depth:** Embedded in agent frameworks
- **Insurance Partnerships:** Exclusive relationships

Go-To-Market Strategy

Phase 1: Foundation (Months 1-12)

Target: AI agent platform providers

- Integrate with top 20 agent frameworks (LangChain, AutoGPT, CrewAI, etc.)
- Launch free tier for developers (100K queries/month)
- Establish data partnerships with cloud providers
- Publish Agent Trust Whitepaper

Key Metrics: - 1,000 platform integrations - 100M monthly score queries - 50 enterprise pilots

Phase 2: Scale (Months 12-24)

Target: Enterprise adopters & insurance industry

- Enterprise sales motion (\$50K+ contracts)
- Insurance carrier partnerships (Lloyd's, AIG, Munich Re)
- Regulatory engagement (SEC, FTC, EU AI Office)
- Agent certification program launch

Key Metrics: - 500 enterprise customers - 3 major insurance partnerships - 1B monthly queries - \$50M ARR

Phase 3: Dominance (Months 24-36)

Target: Industry standard status

- Open standard proposal (Agent Trust Protocol)
- Government/regulatory adoption
- International expansion
- M&A for adjacent capabilities

Key Metrics: - 5,000+ enterprise customers - 50%+ market share in enterprise AI trust - \$300M+ ARR - Regulatory mandates in 3+ jurisdictions

Technology Architecture

System Design

AgentCredit Platform

Identity Registry	Behavioral Observatory	Scoring Engine (ML Pipeline)
----------------------	---------------------------	---------------------------------

Trust Intelligence Graph
(Agent relationships, interaction history, outcomes)

Trust API (Sub-10ms)	Forensics Console	Insurance Gateway (Claims + Pricing)
-------------------------	----------------------	---

Agent Frameworks	Agent Platforms	Agent Insurers
---------------------	--------------------	-------------------

Key Technical Capabilities

- 1. Privacy-Preserving Scoring**
 - Federated learning for distributed computation
 - Differential privacy for aggregate insights
 - Zero-knowledge proofs for verification without exposure
- 2. Real-Time Performance**
 - <10ms API response times
 - 99.99% uptime SLA
 - Global edge deployment
- 3. Fraud Detection**
 - Anomaly detection for compromised agents
 - Sybil attack prevention
 - Behavioral fingerprinting
- 4. Interoperability**
 - REST, GraphQL, gRPC APIs

- OpenTelemetry integration
- Major cloud provider partnerships

Team Requirements

Founding Team (4-6 people)

Role	Profile
CEO	Enterprise SaaS leader, fintech/insurtech background
CTO	Distributed systems, ML infrastructure at scale
Chief Science Officer	ML/AI safety, reputation systems expertise
VP Engineering	Platform infrastructure, API product experience
VP Business Development	Insurance/financial services relationships
General Counsel	AI regulation, data privacy expertise

Advisory Board

- Former FICO executive
- AI safety researcher (MIRI, Anthropic, OpenAI background)
- Insurance industry leader
- AI policy expert (government experience)

Hiring Plan

Timeline	Headcount	Focus Areas
Month 6	25	Engineering, Product, Data Science
Month 12	60	Sales, Customer Success, Partnerships
Month 18	120	International, Enterprise, Compliance
Month 24	200	Scale all functions

Financial Projections

Revenue Forecast

Year	ARR	Growth
Year 1	\$5M	-
Year 2	\$50M	900%
Year 3	\$300M	500%
Year 4	\$800M	167%
Year 5	\$1.5B	88%

Funding Requirements

Round	Amount	Use of Funds	Timeline
Seed	\$8M	Core platform, initial team	Q3 2026
Series A	\$40M	Scale platform, enterprise sales	Q2 2027

Round	Amount	Use of Funds	Timeline
Series B	\$150M	Market dominance, international	Q1 2028
Series C	\$400M	Insurance integration, M&A	Q4 2028

Path to Profitability

- **Break-even:** Month 30
- **EBITDA Positive:** Month 36
- **Target EBITDA Margin:** 35% at scale

Risk Analysis & Mitigation

Risk	Probability	Impact	Mitigation
Major platform builds competing solution	Medium	High	Network effects, insurance integration as moat
Regulatory fragmentation	Medium	Medium	Active policy engagement, flexible architecture
Data privacy challenges	Medium	Medium	Privacy-preserving tech, transparent policies
Slow enterprise adoption	Low	High	Developer-first strategy, free tier
Technical scaling challenges	Low	Medium	Proven infrastructure patterns, strong team

Exit Opportunities

Potential Acquirers

1. **Salesforce** — CRM + AI trust layer (\$15-25B)
2. **Microsoft** — Azure AI governance (\$20-30B)
3. **Visa/Mastercard** — Transaction trust infrastructure (\$10-20B)
4. **AIG/Munich Re** — Insurance infrastructure (\$8-15B)
5. **Moody's/S&P** — Credit rating expansion (\$12-20B)

IPO Path

- Target: 2030-2031
- Valuation: \$20-40B
- Comparable: FICO (\$20B), Moody's (\$70B), Verisk (\$35B)

Why Now?

1. **Inflection Point:** AI agent adoption accelerating exponentially
2. **Trust Gap Widening:** High-profile agent failures creating demand
3. **Regulatory Pressure:** EU AI Act, SEC guidance driving compliance needs

4. **Insurance Demand:** Carriers desperate for underwriting intelligence
 5. **Infrastructure Window:** Standards not yet set — opportunity to define them
-

The Vision

By 2030, every AI agent interaction will be scored by AgentCredit.

Just as FICO became the standard for consumer creditworthiness, AgentCredit will become the universal trust layer for the agentic economy — enabling trillions of dollars in agent-to-agent commerce while protecting humans and organizations from AI risks.

The FICO Score for the Age of AI Agents.

“In the human economy, trust is built over years. In the agent economy, trust must be computed in milliseconds. AgentCredit is that computation.”

Call to Action

We’re raising an \$8M seed round to: - Build the core scoring platform - Integrate with top 20 agent frameworks - Establish insurance partnerships - Hire world-class team

Contact: investors@agentcredit.ai

Generated by The Godfather — May 18, 2026