

AgentOps: The Operating System for Enterprise AI Agents

One-Liner: The infrastructure layer that deploys, monitors, governs, and secures AI agents across the enterprise — the “Kubernetes for autonomous AI.”

Date: May 20, 2026

Drop: Morning

Valuation Target: \$10B+ (Infrastructure Layer)

The Problem

We’re in the middle of an AI agent explosion. Every major company is deploying autonomous AI systems — from customer service bots to code generators to financial analysts. But there’s a critical gap:

No one has built the infrastructure layer for governing AI agents at scale.

Today’s reality: - **Shadow AI Agents:** Employees are spinning up agents without IT oversight - **Zero Observability:** Companies can’t see what their AI agents are doing or why - **Compliance Nightmares:** Regulators are demanding AI audit trails that don’t exist - **Security Black Holes:** Agents with API access are attack vectors with no monitoring - **Resource Chaos:** No way to track compute costs, rate limits, or model usage across agents

The enterprise AI agent market is projected to hit **\$150B by 2028**, but 80% of companies lack any infrastructure to manage what they’re deploying.

The Solution: AgentOps

AgentOps is the operating system for enterprise AI agents — providing the deployment, monitoring, governance, and security infrastructure that every company running AI agents needs.

Core Platform Components

1. Agent Registry & Lifecycle Management

- Central catalog of all AI agents in your organization
- Version control and rollback capabilities
- Deployment pipelines (dev → staging → production)
- Agent “birth certificates” tracking model, prompts, permissions

2. Real-Time Observability

- Trace every decision an agent makes
- Token usage, latency, cost tracking per agent
- Decision audit logs for compliance
- Anomaly detection: “Agent X is behaving differently”

3. Governance & Compliance Layer

- Policy-as-code: Define what agents can/cannot do
- Role-based access control for agent capabilities
- Pre-built compliance templates (SOC 2, GDPR, HIPAA)
- Automated compliance reporting for regulators

4. Security & Guardrails

- Agent sandboxing: Limit blast radius
- Prompt injection detection and prevention
- Credential management for agent API access
- Kill switches and circuit breakers

5. Agent Orchestration

- Multi-agent workflow coordination
 - Inter-agent communication protocols
 - Resource allocation and scheduling
 - Dead letter queues for failed agent tasks
-

Why Now?

1. The Agent Tipping Point (2025-2026)

OpenAI's Operator, Anthropic's computer use, Google's Project Astra — the major labs have all shipped agent frameworks. Adoption is exploding, but governance is an afterthought.

2. Regulatory Pressure Is Mounting

The EU AI Act mandates audit trails for high-risk AI systems. US agencies are drafting similar requirements. Companies need compliance infrastructure yesterday.

3. Enterprise Security Teams Are Panicking

Every AI agent with API access is a potential attack surface. CISOs are demanding visibility and control they don't have.

4. Cost Optimization Demands

Companies are spending millions on AI inference with no visibility into which agents are consuming what. CFOs want metering and chargebacks.

Market Opportunity

Segment	TAM (2028)
AI Agent Orchestration	\$45B
AI Observability & Monitoring	\$28B
AI Governance & Compliance	\$35B
AI Security	\$42B
Total Addressable Market	\$150B

Target Customers

- **Fortune 500 enterprises** deploying 100+ AI agents
- **Financial services** with strict compliance requirements
- **Healthcare organizations** needing HIPAA-compliant AI
- **Tech companies** building AI-native products
- **Government agencies** with security mandates

Product Tiers

AgentOps Community (Free)

- Up to 5 agents
- Basic observability
- 7-day log retention
- Community support

AgentOps Team (\$500/month)

- Up to 50 agents
- Full observability suite
- 90-day log retention
- Compliance templates
- Email support

AgentOps Enterprise (Custom)

- Unlimited agents
 - Custom compliance frameworks
 - Unlimited retention
 - SSO, RBAC, audit logs
 - Dedicated support + SLAs
 - On-prem deployment option
-

Competitive Landscape

Company	Focus	Gap
LangSmith	LLM observability	No governance, limited enterprise features
Weights & Biases	ML experiment tracking	Not built for agents, no compliance
DataDog	General observability	Generic, not AI-native
Snyk	Code security	Not AI-specific
AgentOps	Full agent lifecycle	Unified platform for enterprise AI agents

Our Moat

1. **First-mover in agent governance:** No one owns this category yet
 2. **Data network effects:** More agents = better anomaly detection
 3. **Compliance lock-in:** Once you build compliance on us, switching is painful
 4. **Enterprise distribution:** Fortune 500 relationships are sticky
-

Go-to-Market Strategy

Phase 1: Developer Adoption (Months 1-6)

- Open-source the observability layer

- Free tier for individual developers
- Integration with popular frameworks (LangChain, CrewAI, AutoGen)
- Content marketing: “The State of Enterprise AI Agents” report

Phase 2: Team Expansion (Months 6-12)

- Self-serve paid tier
- Slack/Teams integrations
- SOC 2 compliance certification
- Case studies with early adopters

Phase 3: Enterprise Sales (Months 12-24)

- Dedicated sales team for Fortune 500
- Compliance-led selling (CISO + CCO personas)
- System integrator partnerships (Accenture, Deloitte)
- Industry-specific solutions (FinServ, Healthcare, Gov)

Business Model

Revenue Streams

1. **Subscription SaaS:** Per-agent/per-seat pricing
2. **Usage-based:** Trace volume, log retention
3. **Professional Services:** Custom compliance frameworks, integrations
4. **Marketplace:** Third-party governance policies, connectors

Unit Economics (Target)

- **ACV:** \$150K (enterprise), \$6K (team)
- **Gross Margin:** 80%+
- **CAC Payback:** 12 months
- **Net Revenue Retention:** 140%+

Founding Team Requirements

Role	Profile
CEO	Enterprise SaaS exec, sold to Fortune 500
CTO	Ex-FAANG, distributed systems + ML ops
CPO	Developer tools background, DevRel experience
VP Engineering	Built observability platforms at scale
Head of Security	Ex-enterprise CISO or security vendor

Funding Roadmap

Round	Amount	Use of Funds	Milestones
Seed	\$5M	MVP, 10 design partners	Launch, 50 paying teams
Series A	\$25M	GTM, enterprise features	500 customers, \$5M ARR

Round	Amount	Use of Funds	Milestones
Series B	\$80M	Scale sales, international	50 enterprise logos, \$30M ARR
Series C	\$200M	Category dominance	\$100M ARR, IPO prep

Risk Analysis

Technical Risks

- **Agent fragmentation:** Many agent frameworks, need broad integrations
- **Mitigation:** Start with top 3 frameworks, build abstraction layer

Market Risks

- **Big tech bundles:** AWS/Azure/GCP add agent management
- **Mitigation:** Go deep on compliance/governance where they won't compete

Execution Risks

- **Enterprise sales cycle:** Long and expensive
 - **Mitigation:** Product-led growth to land, expand with sales
-

The Vision

In 5 years, AgentOps will be the control plane for all AI agents in the enterprise — as essential as Kubernetes is for containers or Okta is for identity.

Every autonomous AI system will be registered, monitored, and governed through AgentOps. We'll be the system of record for AI compliance, the dashboard every CISO checks, and the platform every AI team builds on.

The AI agent revolution is here. Someone needs to build the infrastructure to keep it from going off the rails.

That someone is AgentOps.

Key Metrics to Track (Year 1)

- Agents under management: 10,000+
 - Enterprise customers: 25+
 - ARR: \$5M
 - NRR: 140%+
 - GitHub stars (open-source): 5,000+
-

Call to Action

The window to own the AI agent infrastructure category is **18 months**. After that, the market will consolidate or big tech will bundle.

Now is the time to build AgentOps.

Generated by The Godfather / Morning Drop / May 20, 2026